

6-Month Roadmap to Becoming an AI Engineer: A Step-by-Step Guide

Month	Week	Concepts	Goals	Tips
Month 1	Week 1	Python Basics (Variables, Loops, Functions)	Understand Python syntax, build first project	Focus on mastering basics, practice daily
Month 1	Week 2	OOP, NumPy, Pandas	Learn OOP, libraries for AI	Understand Python libraries thoroughly
Month 1	Week 3	Linear Algebra, Calculus	Master linear algebra, calculus for AI	Use visual aids to understand math concepts
Month 1	Week 4	Probability, Statistics	Apply probability, statistics in AI	Practice with real datasets for statistics
Month 2	Week 5	Intro to Machine Learning (Linear Regression)	Build basic linear regression model	Focus on understanding the core of ML algorithms
Month 2	Week 6	Logistic Regression, Classification	Classify data with logistic regression	Work on practical projects, apply classification
Month 2	Week 7	Decision Trees, Random Forest	Build decision tree models, apply Random Forest	Explore hyperparameter tuning
Month 2	Week 8	K-means, Clustering	Cluster data using K-means	Visualise clusters to understand data
Month 3	Week 9	Neural Networks, Backpropagation	Understand how neural networks work	Practice coding small NN from scratch
Month 3	Week 10	Convolutional Neural Networks (CNNs)	Build CNN model for image classification	Experiment with different architectures of CNN
Month 3	Week 11	Recurrent Neural Networks (RNNs)	Use RNNs for sequential data	Understand time-series data
Month 3	Week 12	Autoencoders, GANs, Transfer Learning	Apply advanced DL techniques (GANs, Transfer Learning)	Learn to use pre-trained models effectively
Month 4	Week 13	NLP: Text Preprocessing, Sentiment Analysis	Specialise in NLP, create sentiment analysis tool	Focus on text preprocessing techniques
Month 4	Week 14	Computer Vision: Object Detection, Image Segmentation	Specialise in Computer Vision, build object detection	Work with pre-trained CNN models for quick results
Month 4	Week 15	Reinforcement Learning, Q-Learning	Build an AI agent with reinforcement learning	Test reinforcement learning models in OpenAI Gym

Month 4	Week 16	Portfolio Building and Showcase	Create portfolio, document projects	Highlight most impressive projects in your portfolio
Month 5	Week 17	Model Deployment with Flask, Docker	Deploy AI models on cloud platforms	Use free cloud credits for AWS/Google Cloud
Month 5	Week 18	Deploying on AWS/Google Cloud	Understand and apply cloud deployment techniques	Focus on troubleshooting deployment issues
Month 5	Week 19	MLOps, CI/CD	Implement CI/CD for AI models	Learn version control (Git) for models
Month 5	Week 20	Scaling AI Models with Kubernetes	Learn to scale models for large datasets	Understand Docker and Kubernetes in detail
Month 6	Week 21	Resume and LinkedIn Optimization	Create a professional AI-specific resume	Keep your LinkedIn updated with projects
Month 6	Week 22	Technical Interview Prep (LeetCode, System Design)	Practice coding interviews and AI system design	Simulate mock interviews with a peer
Month 6	Week 23	Networking, AI Conferences	Network and attend AI events	Join online AI communities for networking
Month 6	Week 24	Job Application and Final Interviews	Apply to AI jobs, practice interviews	Tailor resume and cover letter for each job

Table of Contents

1. Introduction: The Power of AI

- What is AI and why it's the future?
- Career opportunities in AI
- Why you should follow this roadmap

2. Phase 1: Fundamentals (Weeks 1-4)

- Python Programming (Weeks 1-2)
 - Why Python for AI?
 - Python basics to master: Variables, Loops, Functions, and OOP

- Key libraries: NumPy and Pandas
 - Project: Data manipulation using Pandas
- Mathematics for AI (Weeks 3-4)
 - Key areas to focus on: Linear Algebra, Calculus, and Probability
 - Tools to simplify math for AI: Wolfram Alpha
 - Project: Linear Algebra with Python (vectors, matrices)

3. Phase 2: Introduction to Machine Learning (Weeks 5-8)

- Basics of Machine Learning (Weeks 5-6)
 - Supervised vs Unsupervised Learning
 - Key ML models: Linear Regression, Decision Trees
 - Project: Predict housing prices using Linear Regression
- Hands-on Projects (Weeks 6-8)
 - Importance of practical projects
 - Project Ideas: Spam classifier using Logistic Regression
 - Tools: Google Colab for training models

4. Phase 3: Deep Learning and Neural Networks (Weeks 9-12)

- Deep Learning Basics (Weeks 9-10)
 - What are Neural Networks?
 - CNNs and RNNs: Concepts and use cases
 - Project: Build a digit recognizer using CNNs
- Frameworks: TensorFlow and PyTorch (Weeks 11-12)
 - TensorFlow vs PyTorch: Which one to choose?
 - Tools: Pre-trained models from TensorFlow Hub
 - Project: Real-time object detection with CNNs

5. Phase 4: Specialization and Portfolio Building (Weeks 13-16)

- Specialization (Weeks 13-14)

- Choosing a specialization: NLP, Computer Vision, or Reinforcement Learning
- Project Ideas: NLP chatbot or Computer Vision real-time object detection

- Portfolio Projects (Weeks 15-16)

- How to create a portfolio that stands out
- Proper documentation, README files, and showcasing projects
- Platforms: GitHub for portfolio building

6. Phase 5: AI Engineering Skills & Deployment (Weeks 17-20)

- Model Deployment (Weeks 17-18)

- Deploying models using Flask or Streamlit
- Scaling models with cloud platforms: AWS, Google Cloud
- Project: Deploy a real-time recommendation system

- Version Control and MLOps (Weeks 19-20)

- Using Git for version control
- Docker for containerization and seamless deployment
- Tools: Docker, Git for managing AI projects

7. Phase 6: Job Preparation & Networking (Weeks 21-24)

- Job Preparation (Weeks 21-22)

- Building a strong AI resume and LinkedIn profile
- Preparing for technical interviews (LeetCode, InterviewBit)

- Networking (Weeks 23-24)

- Importance of networking in AI: Online communities and conferences
- Platforms: LinkedIn, AngelList for job opportunities

8. Free Resources for Learning AI

- Online courses, tutorials, and platforms
- Best resources for mastering Python, Machine Learning, and Deep Learning

9. Conclusion: Your Journey to Becoming an AI Engineer

- How to stay consistent and motivated
 - Final words of advice for a successful career in AI
-

Introduction: The Power of AI

Artificial Intelligence (AI) is revolutionizing the world we live in. From virtual assistants like Siri and Google to self-driving cars, AI is everywhere. But how does AI work, and more importantly, how can you build a rewarding career in AI?

In this guide, we'll walk you through a 6-month roadmap to becoming an AI Engineer, from mastering Python programming and mathematics to building real-world AI projects and deploying them. If you invest 3-4 hours daily in this roadmap, you'll be job-ready with a strong portfolio of AI projects.

Phase 1: Fundamentals (Weeks 1-4)

The first phase focuses on building a strong foundation in **Python Programming, Mathematics for AI, and Data Structures & Algorithms (DSA)**. By the end of this phase, you should be comfortable with Python basics, AI-related math, and essential data structures.

Week 1: Python Basics

Day 1-2: Introduction to Python Programming

- What is Python? Why is it used in AI?
- Installing Python, setting up your environment (Jupyter, VS Code)
- Python syntax, variables, data types (int, float, str, bool)

Day 3: Operators in Python

- Arithmetic, logical, and comparison operators
- Assignment and bitwise operators

Day 4-5: Control Flow

- If-else statements, nested conditions
- Loops: For and while loops, nested loops

Day 6: Functions in Python

- Defining functions, function parameters, return statements
- Global and local scope, using `*args` and `**kwargs`

Day 7: Weekend Project

- **Project:** Build a simple temperature converter that converts Celsius to Fahrenheit and vice versa. Incorporate input validation and function reusability.
-

Week 2: Advanced Python Concepts & Libraries

Day 1-2: Object-Oriented Programming (OOP)

- Classes and objects, constructors
- Inheritance, polymorphism, encapsulation

Day 3-4: Python Libraries for AI

- NumPy: Arrays, reshaping, and slicing
- Pandas: DataFrames, basic data manipulation

Day 5: File Handling

- Reading and writing files in Python
- Working with CSV files (using Pandas)

Day 6: Exception Handling

- Try-except blocks, custom exceptions
- Best practices for error handling in AI

Day 7: Weekend Project

- Project: Create a basic data analysis script using Pandas. Load a CSV file of a dataset (like weather data) and compute basic statistics (mean, median, etc.).
-

Week 3: Mathematics for AI – Linear Algebra & Calculus

Day 1-2: Introduction to Linear Algebra

- Vectors and scalars
- Matrix operations (addition, multiplication)

Day 3: Eigenvalues and Eigenvectors

- Definitions and applications in AI
- Diagonalization of matrices

Day 4-5: Calculus for AI

- Derivatives and gradients (understand how models learn)
- Chain rule and its importance in backpropagation

Day 6: Partial Derivatives

- Using partial derivatives in optimization problems

Day 7: Weekend Project

- **Project:** Implement matrix multiplication from scratch in Python. Use NumPy to verify results. This project will help you understand how linear algebra is used in deep learning.
-

Week 4: Probability & Statistics for AI

Day 1-2: Basic Probability Concepts

- Probability distributions (Normal, Poisson, Binomial)
- Bayes Theorem

Day 3-4: Statistics for AI

- Mean, median, mode, variance, and standard deviation
- Hypothesis testing

Day 5-6: Probability in Machine Learning

- Overfitting and underfitting, bias-variance trade-off

Day 7: Weekend Project

- **Project:** Create a Python script that simulates flipping a coin 1000 times. Analyze the results using statistical measures (mean, variance) and create visualizations using Matplotlib.
-

Phase 2: Introduction to Machine Learning (Weeks 5-8)

This phase dives into machine learning fundamentals. By the end, you'll have built your first machine learning models and understand how to process data.

Week 5: Introduction to Machine Learning (ML)

Day 1-2: What is Machine Learning?

- Supervised vs unsupervised learning
- Training, validation, and test sets

Day 3-4: Linear Regression

- Understanding regression problems
- Implementing Linear Regression in Python using Scikit-learn

Day 5: Cost Function and Gradient Descent

- Mean Squared Error (MSE)
- Introduction to optimization: Gradient Descent

Day 6: Evaluation Metrics

- R-squared, Adjusted R-squared

Day 7: Weekend Project

- **Project:** Build a Linear Regression model to predict house prices using a dataset (like the Boston Housing Dataset). Evaluate the model using R-squared.
-

Week 6: Logistic Regression and Classification

Day 1-2: Introduction to Classification

- Logistic Regression for binary classification
- Sigmoid function

Day 3-4: Confusion Matrix

- Precision, recall, F1 score, and accuracy

Day 5: Regularization Techniques

- L1 and L2 regularization to prevent overfitting

Day 6: Multiclass Classification

- One-vs-rest and One-vs-one strategies

Day 7: Weekend Project

- **Project:** Build a Logistic Regression model to classify spam emails. Use precision, recall, and F1 score to evaluate your model.
-

Week 7: Decision Trees and Ensemble Methods

Day 1-2: Introduction to Decision Trees

- How decision trees work
- Gini index, entropy, and information gain

Day 3-4: Random Forest

- Bagging and Random Forest algorithm
- Tuning hyperparameters

Day 5: Gradient Boosting

- Boosting, Gradient Boosting, and AdaBoost

Day 6: XGBoost

- Introduction to XGBoost and feature importance

Day 7: Weekend Project

- Project: Build a Random Forest model to classify bank customers as likely or unlikely to leave based on transaction data. Use XGBoost to improve the model performance.
-

Week 8: Unsupervised Learning

Day 1-2: Introduction to Clustering

- K-means clustering
- Elbow method for choosing K

Day 3-4: Hierarchical Clustering

- Agglomerative and divisive clustering

- Dendograms

Day 5-6: Dimensionality Reduction

- Principal Component Analysis (PCA)

Day 7: Weekend Project

- **Project:** Use K-means clustering on customer segmentation data to group similar customers based on purchasing behavior.
-

Phase 3: Deep Learning and Neural Networks (Weeks 9-12)

This phase covers Deep Learning concepts and how to build neural networks.

Week 9: Introduction to Neural Networks

Day 1-2: What is a Neural Network?

- Perceptrons and activation functions
- Structure of a neural network: input, hidden, and output layers

Day 3-4: Backpropagation and Optimization

- How backpropagation works
- Gradient descent for training neural networks

Day 5: Activation Functions

- Sigmoid, ReLU, Tanh, and Softmax

Day 6: Loss Functions

- Cross-entropy loss, Mean Squared Error

Day 7: Weekend Project

- **Project:** Build a simple feedforward neural network from scratch in Python to classify handwritten digits from the MNIST dataset.
-

Week 10: Convolutional Neural Networks (CNNs)

Day 1-2: Introduction to CNNs

- How CNNs work (convolutions, pooling)
- CNN architecture

Day 3-4: Layers in CNNs

- Convolutional layers, pooling layers, fully connected layers

Day 5-6: Regularization in CNNs

- Dropout, batch normalization

Day 7: Weekend Project

- **Project:** Build a CNN to classify CIFAR-10 images (like cars, cats, dogs, etc.).
-

Week 11: Recurrent Neural Networks (RNNs)

Day 1-2: Introduction to RNNs

- Sequential data and time-series prediction
- How RNNs process sequential information

Day 3-4: Long Short-Term Memory (LSTM)

- LSTM vs vanilla RNNs
- Applications of LSTM (e.g., speech recognition, language modeling)

Day 5-6: Gated Recurrent Units (GRU)

- GRU vs LSTM

Day 7: Weekend Project

- Project: Build an LSTM model to predict stock prices based on historical data.
-

Week 12: Advanced Deep Learning

Day 1-2: Autoencoders

- What are autoencoders?
- Applications of autoencoders (e.g., image denoising)

Day 3-4: Generative Adversarial Networks (GANs)

- Introduction to GANs
- How GANs work (generator and discriminator)

Day 5-6: Transfer Learning

- Using pre-trained models for custom tasks

Day 7: Weekend Project

- **Project:** Use a pre-trained CNN model (e.g., VGG16) for image classification on a custom dataset.
-

Phase 4: Specialization and Portfolio Building (Weeks 13-16)

At this stage, you'll choose a specialization and build real-world projects for your portfolio.

Week 13: Specialization – Natural Language Processing (NLP)

Day 1-2: Introduction to NLP

- What is NLP and why is it important?
- Text preprocessing: tokenization, stop words, stemming, and lemmatization

Day 3-4: Bag of Words and TF-IDF

- How to represent text data for machine learning models
- Implementing Bag of Words and TF-IDF from scratch

Day 5-6: Sentiment Analysis

- Understanding sentiment analysis in text data
- Building a basic sentiment classifier using logistic regression

Day 7: Weekend Project

- **Project:** Build a sentiment analysis tool to classify movie reviews as positive or negative using the IMDB dataset. Use TF-IDF for text representation and a machine learning model (Logistic Regression, Naive Bayes) to classify reviews.
-

Week 14: Specialization - Computer Vision

Day 1-2: Image Preprocessing

- Image data augmentation: cropping, scaling, rotation
- Converting images to arrays, normalization, and preprocessing

Day 3-4: Object Detection

- Understanding object detection techniques
- Using pre-trained CNN models for object detection (YOLO, SSD)

Day 5-6: Image Segmentation

- What is image segmentation?
- Implementing image segmentation using U-Net architecture

Day 7: Weekend Project

- **Project:** Build an object detection system that identifies objects in real-time using a webcam. Use a pre-trained model like YOLO and implement the system using OpenCV.
-

Week 15: Specialization - Reinforcement Learning

Day 1-2: Introduction to Reinforcement Learning

- Key concepts: agents, environments, rewards, and actions
- Exploration vs. exploitation in reinforcement learning

Day 3-4: Q-Learning Algorithm

- Understanding Q-Learning and its applications
- Implementing Q-Learning from scratch

Day 5-6: Deep Q-Networks (DQN)

- Introduction to DQNs: Using deep learning with Q-learning
- Applications in AI gaming and robotics

Day 7: Weekend Project

- Project: Use OpenAI Gym to train an AI agent that plays a simple game (like CartPole). Implement Q-Learning to optimize the agent's behavior.
-

Week 16: Portfolio Projects and Showcase

Day 1-3: Review and Clean Up Portfolio Projects

- Go back to all your projects, clean up code, add detailed comments and documentation
- Ensure your GitHub repository is well-organized with proper README files

Day 4-5: Creating a Portfolio Website

- Use simple HTML/CSS or platforms like GitHub Pages to create a personal portfolio website where you showcase all your AI projects

Day 6: Writing a Blog Post on Your AI Journey

- Write a detailed blog post on Medium or LinkedIn about your AI journey, showcasing the projects you've built and the skills you've acquired

Day 7: Weekend Project

- **Project:** Create a demo video for one of your key projects (e.g., sentiment analysis or object detection). Walk through the project, explain its significance, and share it on social media or your portfolio site.
-

Phase 5: AI Engineering Skills & Deployment (Weeks 17-20)

This phase focuses on turning your models into real-world products by learning deployment strategies and how to manage AI projects in production environments.

Week 17: Model Deployment Basics

Day 1-2: Introduction to Model Deployment

- Why deployment is important
- Setting up a Flask API for your machine learning model

Day 3-4: Flask + Heroku

- Deploying a simple model using Flask and pushing it to Heroku (free cloud platform)

Day 5-6: Docker Basics

- What is Docker? Why containerization is important for AI models
- Creating a Docker container for your AI project

Day 7: Weekend Project

- Project: Deploy a pre-built machine learning model (e.g., your Logistic Regression spam classifier) to Heroku using Flask. Containerize the model using Docker.
-

Week 18: Advanced Model Deployment

Day 1-2: Deployment on AWS

- Setting up an AWS account and using S3 for data storage
- Deploying your AI model on AWS EC2

Day 3-4: Using AWS Lambda for Serverless Deployment

- What is serverless architecture?
- Deploying a model using AWS Lambda and API Gateway

Day 5-6: Google Cloud Deployment

- Deploying machine learning models on Google Cloud AI Platform
- Using Google Cloud Storage and Cloud Functions

Day 7: Weekend Project

- **Project:** Deploy a CNN-based image classification model (from Week 10) on AWS or Google Cloud and create a web app where users can upload images and get predictions in real-time.
-

Week 19: MLOps - Version Control and Continuous Deployment

Day 1-2: Introduction to MLOps

- What is MLOps? Why it's crucial for AI Engineers?
- Version control for ML models using Git

Day 3-4: CI/CD Pipelines for Machine Learning

- Introduction to Continuous Integration and Continuous Deployment (CI/CD)
- Tools: Jenkins, GitHub Actions for automating model deployment

Day 5-6: Monitoring and Logging

- How to monitor machine learning models in production
- Setting up logging for errors, model performance, and scaling

Day 7: Weekend Project

- **Project:** Set up a simple CI/CD pipeline using GitHub Actions for continuous deployment of your Flask model to Heroku. Implement logging to monitor errors in real-time.
-

Week 20: Scaling AI Systems

Day 1-2: Introduction to Scaling

- How to scale AI systems to handle large traffic
- Load balancing and auto-scaling on AWS and Google Cloud

Day 3-4: Distributed Training

- Using TensorFlow or PyTorch for distributed model training
- Implementing distributed data parallelism

Day 5-6: Using Kubernetes for AI

- What is Kubernetes? How to deploy and manage AI models at scale using Kubernetes

Day 7: Weekend Project

- **Project:** Set up a distributed training environment for your CNN model using TensorFlow and Kubernetes. Train the model on multiple machines and deploy the final model to AWS.
-

Phase 6: Job Preparation & Networking (Weeks 21-24)

In this final phase, you will prepare for job applications, technical interviews, and networking to break into the AI industry.

Week 21: Resume and LinkedIn Optimization

Day 1-2: Building an AI-Specific Resume

- How to highlight AI projects and skills in your resume
- Tailoring your resume for AI roles

Day 3-4: LinkedIn Optimization

- Optimizing your LinkedIn profile for AI-related jobs
- Building a personal brand: Sharing blog posts, projects, and engaging with AI content

Day 5-6: GitHub Optimization

- Ensure your GitHub profile is polished and easy to navigate
- Add detailed READMEs and documentation for your AI projects

Day 7: Weekend Task

- Task: Share a post on LinkedIn showcasing one of your AI projects.
Write a detailed explanation of the project, the challenges faced, and its applications.
-

Week 22: Technical Interview Preparation

Day 1-2: AI Technical Interviews Overview

- What to expect in AI technical interviews
- Most common questions related to machine learning, deep learning, and deployment

Day 3-4: LeetCode for AI Engineers

- Practicing machine learning-related coding problems on LeetCode

Day 5-6: System Design for AI

- How to approach AI system design questions in interviews

- Practice designing AI-based recommendation systems or chatbots

Day 7: Weekend Task

- Task: Simulate a mock technical interview with a friend or online coding platforms. Focus on explaining your AI project in detail.
-

Week 23: Networking and Conferences

Day 1-2: Networking in the AI Industry

- How to connect with AI professionals on LinkedIn
- Attending AI meetups, conferences, and webinars

Day 3-4: Building a Referral Network

- How to get referrals for AI jobs
- The importance of engaging with AI-related online communities (Reddit, Discord)

Day 5-6: AI Conferences

- Explore top AI conferences like AI Expo, Google's AI Summit
- Write a summary of what you learned and post it on LinkedIn

Day 7: Weekend Task

- Task: Attend an online AI webinar or conference. Network with professionals and share your learning experience on LinkedIn.
-

Week 24: Final Job Application and Interview Practice

Day 1-2: Applying for AI Jobs

- Research AI companies and apply for jobs that match your skills
- Tailor your resume and cover letter to each company

Day 3-4: Final Technical Interview Practice

- Practice explaining your portfolio projects in technical interviews
- Participate in coding challenges and mock interviews

Day 5-6: Preparing for HR Interviews

- How to answer behavioral questions
- Discussing your AI journey, challenges, and growth

Day 7: Weekend Task

- **Task:** Apply to at least 10 AI-related jobs, focusing on your resume, cover letter, and portfolio.
-

Conclusion

After following this 6-month roadmap, you will have the necessary skills and practical experience to enter the AI industry. You'll have built a solid portfolio of AI projects, learned how to deploy models, and gained the confidence to tackle AI interviews.

Your journey to becoming an AI Engineer starts now – stay consistent, build projects, and never stop learning!

